


# ***Active Vision***

# Natural Intelligence



- **Attention** is a resource
- Our brain **guides attention** through **eye saccades** (movements)
- We efficiently **sample** the environment and make quick **decisions**

# Artificial Intelligence

- Emulate the eye with a **sensor** like: **RGB camera**, **depth camera**, or both (**RGB-D camera**)
- Make our sensor **mobile**
- **Maximize *Information Gain*** (with *minimal effort*)
- **Optimize** sampling **Efficiency**
- **Controll** the point of view of the mobile sensor via a clever **algorithm**



Replace the **eye** with a **camera**

# Nonmyopic View Planning

- **Detect and Classify** objects in the scene
- Estimate their **pose**

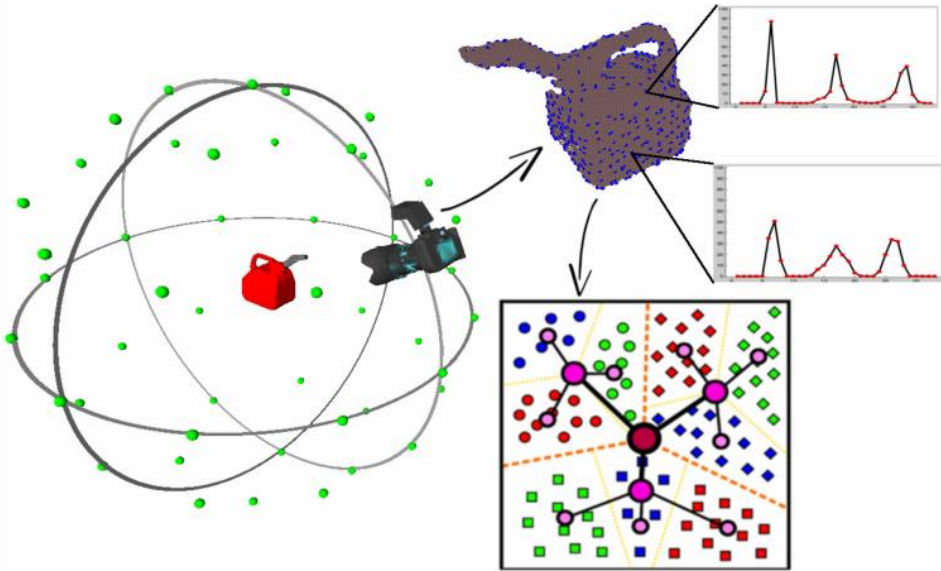
Bridge the gap



## Robot Grasping

- Use a **planner** to estimate a **sequence** of camera poses
- Model using a **POMDP** : **P**artially **O**bservable **M**arkov **D**ecision **P**rocess
- Maximize **long term reward**, i.e. **NOT** Greedy or Short-Sighted/Myopic

# Viewpoint-pose Tree



Features used to construct **VP-Tree**



**No segmenter** component available  
(=> Clutter Problem)



# Myopic approach


- **Greedy**/Short-Sighted/Next-Best-View approach
- Care about **short term reward**/one step horizon
- Can be modeled in many ways i.e. MDP, POMDP, **Neural Nets** etc.



*Especially useful for quadrotor, drone & various flying-robot applications*

# Oracle Neural Net

- Many Greedy approaches still use a POMDP
- This approach uses a trained **Convolutional Neural Network** that approximates an **Oracle Function**

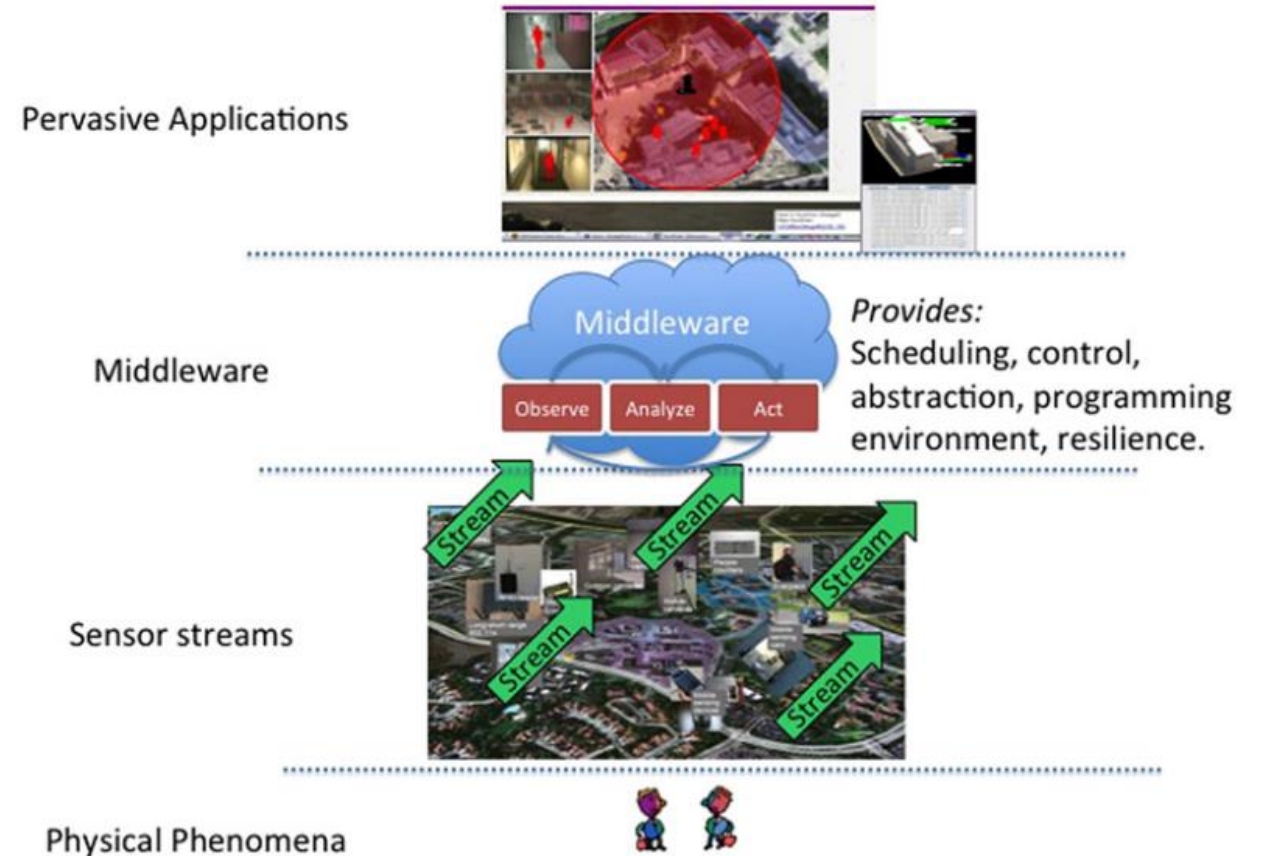


$$\mathcal{L}(X, Y; \theta) = \sum_{i=1}^N \|f(X_i) - Y_i\|_2^2 + \lambda \|\theta\|_2^2$$

- Computes **Next Best View** very fast (inference time of the net is relatively small as compared to a POMDP)
- **No free lunch !** Pay with a **long (offline) training time** and lots of **ground truth** data

# Framework for Monitoring Spaces


- Employ an **approximate POMDP** policy/solution
- Use a **large scale** camera **network** (up to 100 cameras)
- Actuate the pan-**ZOOM**-tilt parameters of cameras




Multiple layers view of **sentient** systems




# Tips and Tricks



Non-overlapping Fields of View




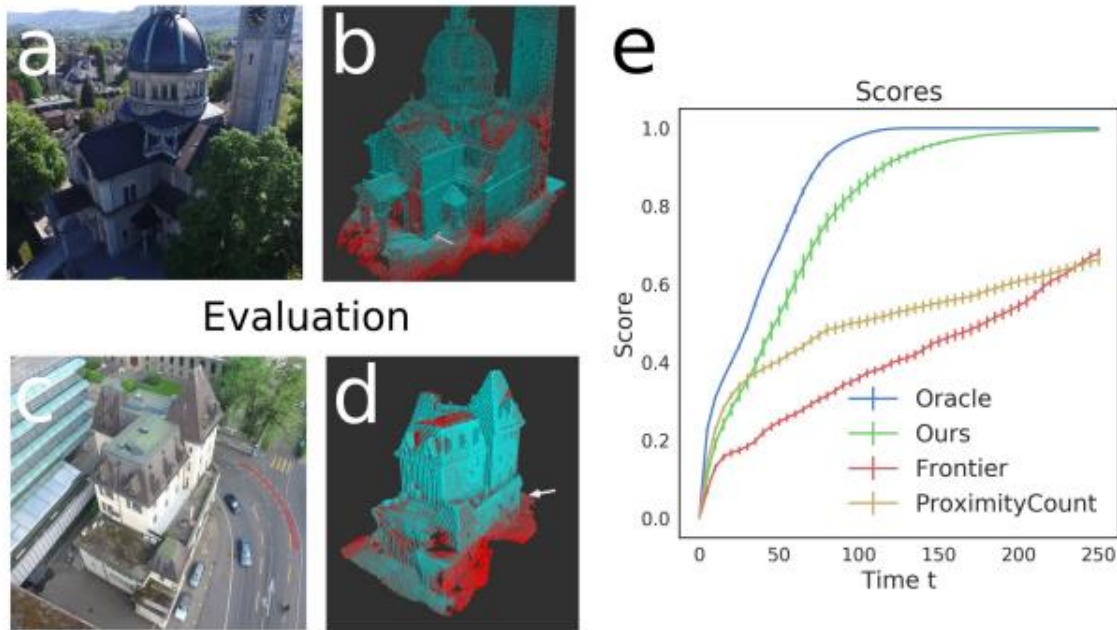
Extract (correlation-) semantics



# Evaluation and Results

		True Hypothesis							Avg Number of Measurements	Avg Movement Cost	Avg Decision Cost	Avg Total Cost
		H(0°)	H(60°)	H(120°)	H(180°)	H(240°)	H(300°)	H( $\emptyset$ )				
Predicted (%)	H(0°)	87.5	2.5	5.0	0.0	0.0	0.0	5.0	2.53	2.81	9.38	14.72
	H(60°)	2.5	80.0	0.0	0.0	0.0	0.0	17.5	2.66	2.52	15.00	20.18
	H(120°)	7.5	0.0	72.5	0.0	0.0	0.0	20.0	3.16	3.43	20.63	27.22
	H(180°)	0.0	0.0	0.0	70.0	10.0	2.5	17.5	2.20	1.72	22.5	26.42
	H(240°)	0.0	0.0	0.0	2.5	75.0	2.5	20.0	2.39	2.51	18.75	23.65
	H(300°)	0.0	0.0	0.0	0.0	5.0	72.5	22.5	2.57	2.18	20.63	25.38
	H( $\emptyset$ )	0.0	0.0	0.97	0.0	0.0	0.97	98.05	2.17	1.93	1.46	5.56
	Overall Average Total Cost:											20.45

## NVP (Nonmyopic View Planner) Approach (on PR2 robot)



Latency of scheduling as a function of number of cameras

## Approximate POMDP Approach

# Concluding Remarks

- 1) **NVP** approach works well in simulation and we can **transfer the encoded knowledge** to a real robot
  - It is not much better quantitatively than a state of the art **Greedy Approach**, but the main advantage is the **adaptive stopping criterion**
  - Does not scale well to **large landscapes**, since solving POMDPs is **exponentially complex**
  - Would benefit from a **custom segmenter** to avoid unnecessary camera movements
- 2) **CNN** approach works much better for **large scenes**
- 3) We can still leverage POMDPs for large scale systems by creatively **engineering clever setups**, employing **simplifying assumptions** and reasonable **approximations**

***“Always ask questions for an  
Active Learning”***

**Active**  
**LEARNING**